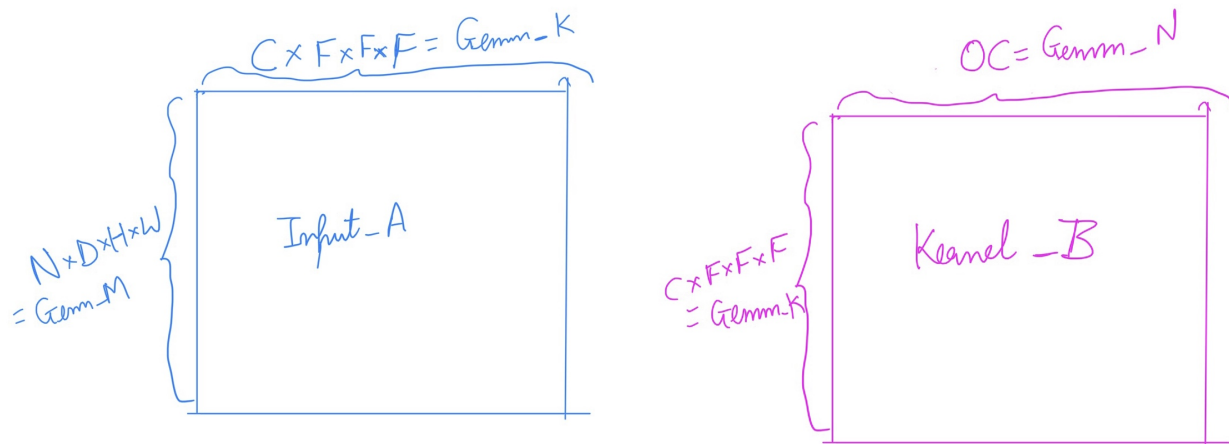


CUTLASS based 3D conv

https://github.com/NVIDIA/cutlass/blob/master/media/docs/implicit_gemm_convolution.md

Adapting optimization to 3D CNN:

Gemm formulation



Initial convolution dimensions:

Output = $N \times D \times H \times W \times OC$

Kernel = $C \times F \times F \times F \times OC$

Data layout = NDHWC

Corresponding gemm dimensions:

A Input: $\text{GEMM_M} \times \text{GEMM_K}$

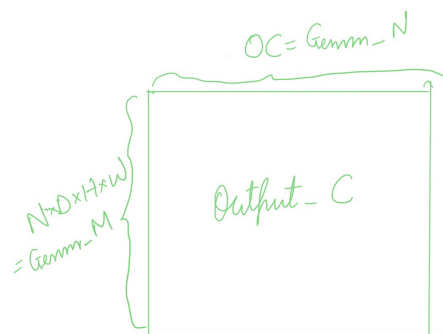
B Kernel: $\text{GEMM_K} \times \text{GEMM_N}$

C Output = $\text{GEMM_M} \times \text{GEMM_N}$

$\text{GEMM_M} = N \times D \times H \times W \times C$

$\text{GEMM_K} = C \times F \times F \times F$

$\text{GEMM_N} = OC$



Implicit gemm indexing

```
for gemm_i in 0:batchsize*outdepth*outheight*outwidth # GEMM_M loop
    for gemm_j in 0:outchannels # GEMM_N loop

        n = gemm_i//tmp_dhw # Indexing the batch size
        nopq_residual = gemm_i % tmp_dhw
        o = nopq_residual//tmp_hw # Indexing the output depth
        opq_residual = nopq_residual%tmp_hw
        p = opq_residual//outwidth # Indexing the output height.
        q = opq_residual%outwidth # Indexing the output width

        accum = 0
        for gemm_k in 0:inchannels*kdim*kdim*kdim # GEMM_K loop

            c = gemm_k//tmp_kdim3 # Indexing input, kernel channel
            ctrs_residual = gemm_k%tmp_kdim3
            t = ctrs_residual//tmp_kdim2
            trs_residual = ctrs_residual%tmp_kdim2
            r = trs_residual//kdim
            s = trs_residual%kdim
            d = o + t # Indexing the input depth e(o, t)
            h = p + r # Indexing the input height f(p, r)
            w = q + s # Indexing the input width g(q, s)
            accum = accum + imgemm_input[n,d,h,w,c]*imgemm_kernel[gemm_j,t,r,s,c]

        imgemm_output[n,o,p,q,gemm_j] = accum
```

Each index `gemm_i` in `GEMM_M` dimension corresponds to a unique (N,O,P,Q) index of the output tensor.

```
n = gemm_i//tmp_dhw # Indexing the batch size
nopq_residual = gemm_i % tmp_dhw
o = nopq_residual//tmp_hw # Indexing the output depth
opq_residual = nopq_residual%tmp_hw
p = opq_residual//outwidth # Indexing the output height.
q = opq_residual%outwidth # Indexing the output width
```

The algorithm partitions the `GEMM_K` dimension into threadblock tiles.

Assigns each threadblock tile to one filter position (t,r,s) and an interval of channels.

Mapping to CUTLASS Gemm matrix multiplication loop

```

#for cta_n in 0:GEMM_N:CTAtileN
for cta_n in 0:outchannels #GEMM_N, gemm_j
  #for cta_m in 0:GEMM_M:CTAtileM
  for cta_m in 0:batchsize*outdepth*outheight*outwidth:CTAtileM #GEMM_M, gemm_i
    #for cta_k in 0:GEMM_K:CTAtileK # Main loop
    for cta_k in 0:inchannels*kdim*kdim*kdim:CTAtileK #GEMM_K, gemm_k

      for warp_n in 0:CTAtileN:WARPtileN
        for warp_m in 0:CTAtileM:WARPtileM
          for warp_k in 0:CTAtileK:WARPtileK # Partitioning the input channels

            for mma_k in 0:WARPtileK
              for mma_n in 0:WARPtileN
                for mma_m in 0:WARPtileM
                  # Execute instruction for accumulating the product

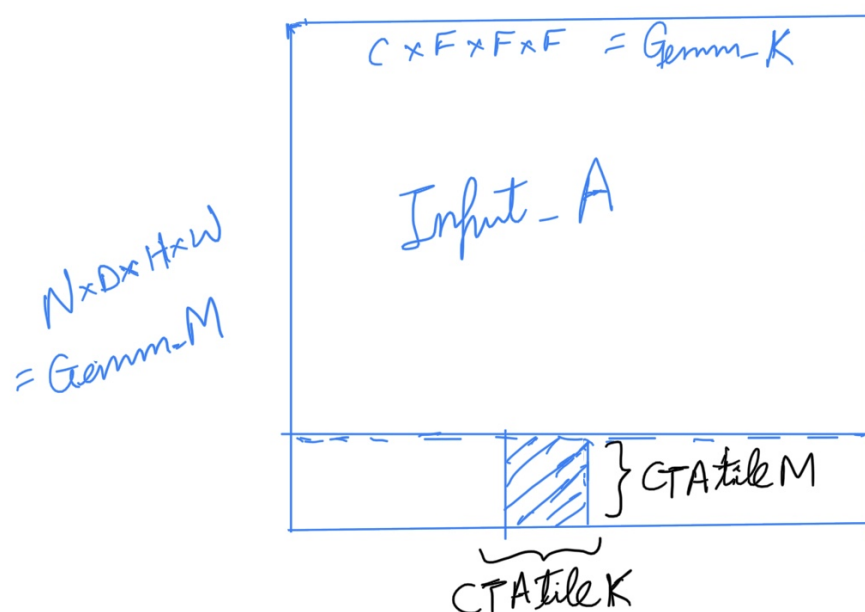
```

One particular iteration of main loop, i.e. when cta_n, cta_m, cta_k are fixed.

The data format for input and kernel is **NDHWC**. So accessing adjacent 'channel index' for a fixed depth, height, width is faster both for input and the filter.

Input computation partitioning

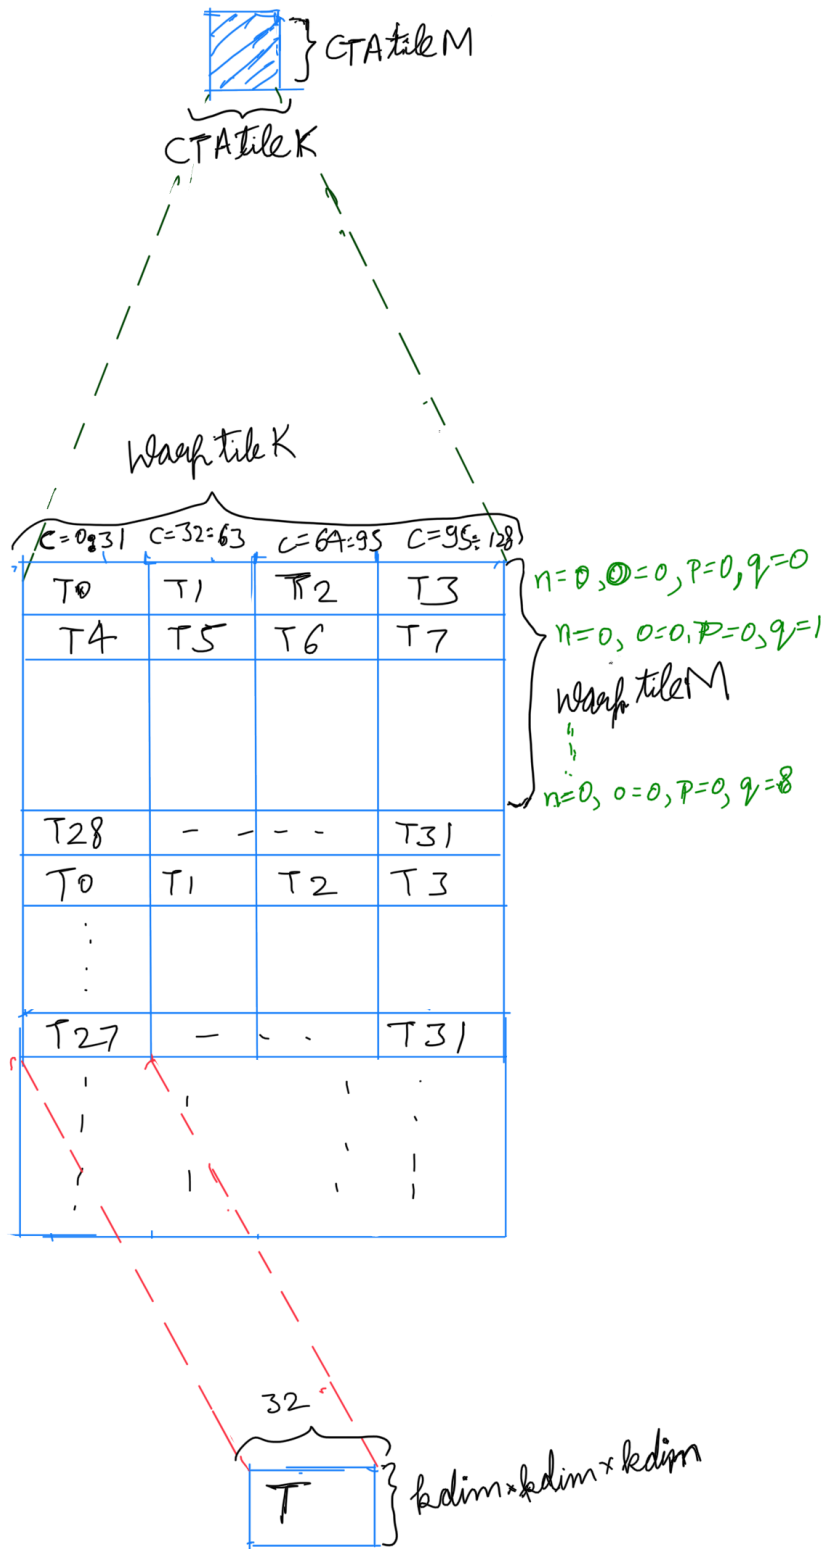
CTAtileM X CTAtileK



- Gemm_K decides the r, s, t coordinates i.e. the indexing of the kernels.
- Each block has a CTAtileM X CTAtileK size input.
- Because the data layout is NDHWC, adjacent channels are accessed by adjacent threads.

WarptileK X WarptileM

- Index along the CTAtileM decides which output values are computed.
- r, s, t index decides which CTA tile or block id.
- A range of channels decided = CTAtileK are used in this block
- Assuming CTAtileK = 128

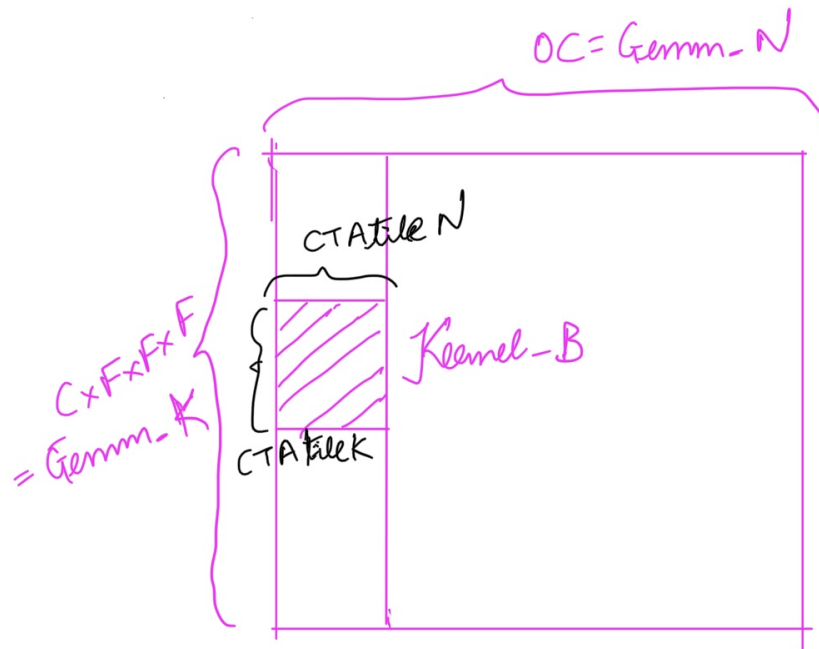


- 32 threads in 1 warp, so $\text{WARPTileK} \times \text{WarptileM} = 32$
- all 4 threads in a row work on 1 output value (indexed by oc, o, p, q)

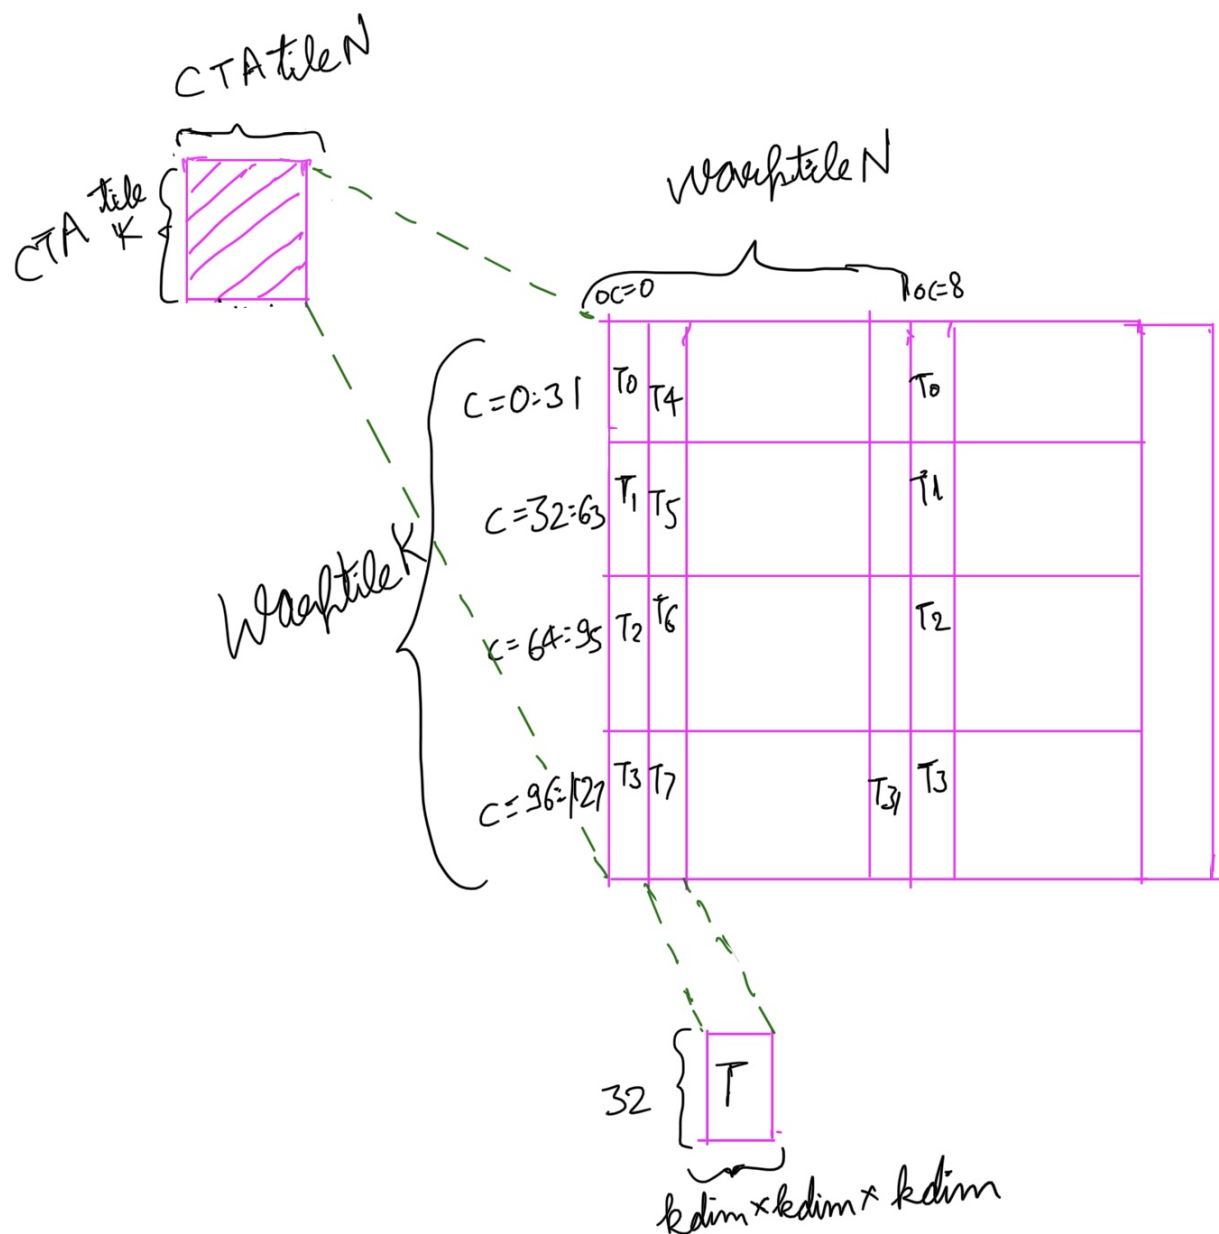
Filter computation partitioning

CTATileK X CTATileN

- Gemm_N is for indexing output channel
- The CTATile selected has a fixed r, s, t based on the block ID
- The range of input channels each block uses is decided by CTATileK
- The range of output channels each block computes is decided by CTATileN



WarptileK X WarptileM



- 32 threads in 1 warp, so $WARPTileK \times WarptileM = 32$
- all 4 threads in a column work on 1 output value indexed by (oc, o, p, q)

Shared memory usage

- The CTAtiles are loaded by each block
- So the CTAtiles corresponding to the input, filter and output should fit inside shared memory

Input = CTATileM * CTATileK

Output = CTATileM * CTATileN

Kernel = CTATileK * CTATileN

Total shared memory needed per block = $[\text{CTATileM} * \text{CTATileK} + \text{CTATileM} * \text{CTATileN} + \text{CTATileK} * \text{CTATileN}]$

this should be less than maximum available CUDA shared memory.

Deciding tile sizes

Goal: maximize the parallelism

Number of blocks launched = $(\text{GEMM_M} / \text{CTATileM}) * (\text{GEMM_N} / \text{CTATileN}) * (\text{GEMM_K} / \text{CTATileK})$

Use soap analysis ??